

Understanding SVM Kernels: A Practical Guide to Choosing the Right Kernel for Your Data

GitHub Repository: [Link](#)

Introduction: Why This Tutorial?

We have probably heard that Support Vector Machines (SVMs) are more popular in machine learning due to the concept of a single word, which is kernels. Kernels make SVMs; they enable classification of complex and high-dimensional data in simpler forms easier. They serve as a potent mathematical device that allows the SVM algorithm to exceed linear separability, which is an important feature to have in dealing with real-life, non-linear problems. The lack of kernels would force SVMs to have linear decision boundaries only, and in most cases, they would be less versatile and effective. However, with the kernels, SVMs are able to process complex patterns and relationships in the data to extent that they are a popular option in a great many classification problems.

Support Vector Machines (SVMs) can solve both simple and complex classification problems on the same basic algorithm with the help of kernels. Nevertheless, using the improper kernel may greatly affect the performance, causing poor outcomes.

There are three most common used kernels of SVM, namely Linear, Polynomial, and Radial Basis Function (RBF), which we will discuss in this tutorial. We shall investigate operations of each kernel, when each should be used, and in what way, to tune the parameters of the kernels. At the end, you will have the knowledge to make sound choices of the correct kernel to use in your machine learning projects and adjust them accordingly.

We are going to use Python and the popular Iris dataset in order to produce explicit visuals that will illustrate the mechanics of how the kernel works. Everything is in the GitHub repository, which you can follow and experiment with on your own.

What You Need to Know: SVM Basics

A Support Vector Machine (SVM) is designed to determine the most suitable line that classifies different categories within a particular data set. As an example, consider when the graph is filled in with data points of two different classes, blue and green, as seen in Figure 1, the SVM algorithm will draw a line (or curve) that approximately separates the two, and at the same time maximizes the distance between the two i.e., the margin. The margin is a very important feature because

the higher the margin the better the model can extrapolate to the unseen data (Cortes and Vapnik, 1995).

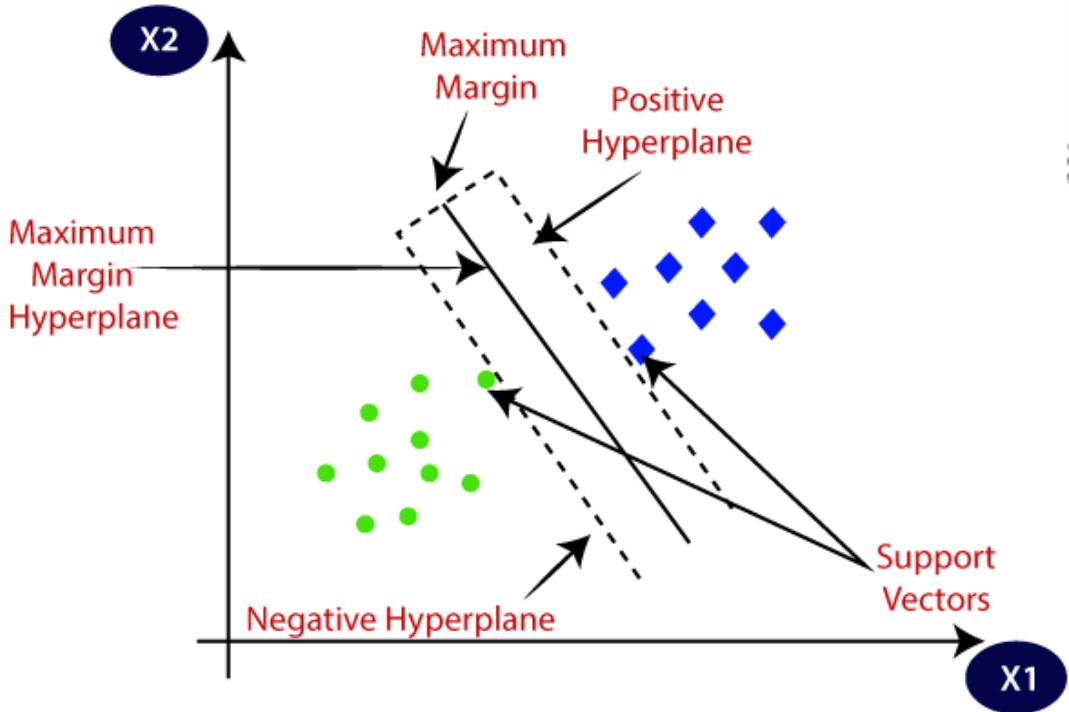


Figure 1: Support Vector Machine (SVM) | Source: [Medium](#)

The large difference between SVM and other algorithms is that it addresses those points in data that are nearest to decision boundary, and these points are called support vectors, as seen in Figure 1. These support vectors are very important in determining the boundary, whilst the others are actually neglected. This property enables SVM to be computationally efficient and less susceptible to overfitting as a result of outliers.

The notion of kernels is specifically important when it comes to real-life data, which is frequently difficult and non-linearly separable. An ordinary linear boundary is not adequate in such circumstances. The kernel trick comes in to face this challenge by allowing the SVM to form non-linear decision boundaries. This is done by mathematically computing the transformation in which the data is converted to a higher-dimensional space without actually taking transformation (Schölkopf et al., 1998).

To illustrate this, imagine one is trying to separate the data points on a flat surface, but a straight line would not work. Assuming that, one could raise some of the points into a three-dimensional space, it may be possible to draw a line that effectively separates them as seen in Figure 2. This transformation is automated by kernels, enabling SVM to deal with a far more important range

of relationships in the data and, in this way, better its classification of non-linearly separable patterns.

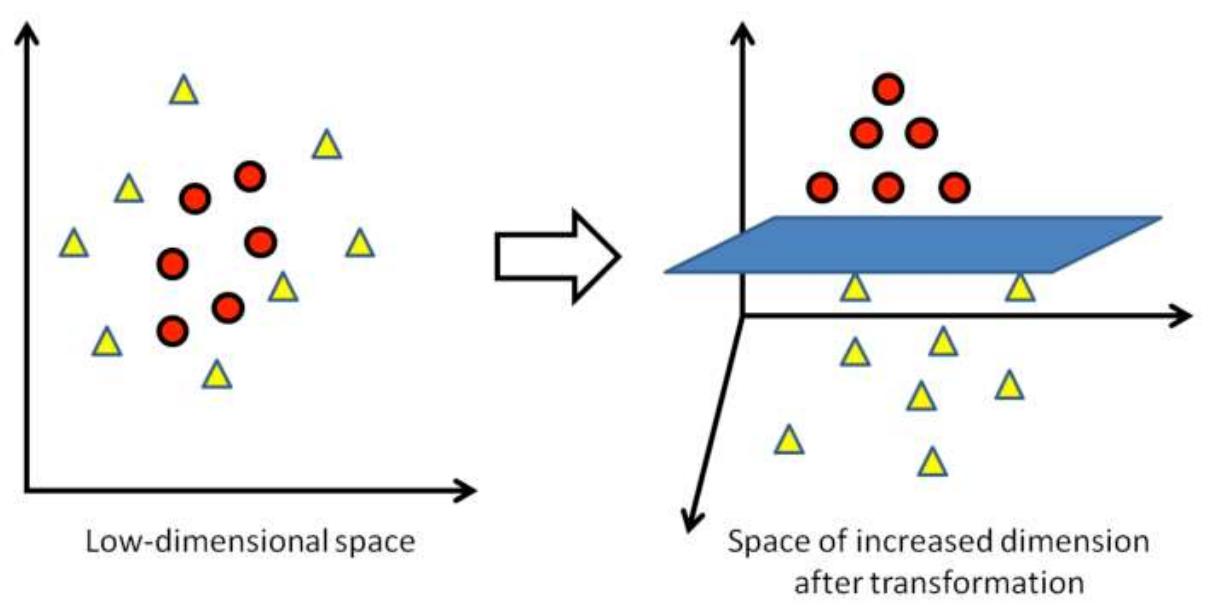


Figure 2: Kernel Trick | Source: [Research Gate](#)

The Dataset: Iris Flowers

Using the famous Iris dataset for this tutorial because it's perfect for demonstrating kernel behavior:

- **150 samples** of iris flowers
- **3 species**: Setosa, Versicolor, and Virginica
- **4 features**: sepal length/width and petal length/width
- **Balanced classes**: 50 samples of each species

For visualization, we will only focus on just two features: petal length and petal width, since they show the clearest separation, and we can then easily plot them in 2D.

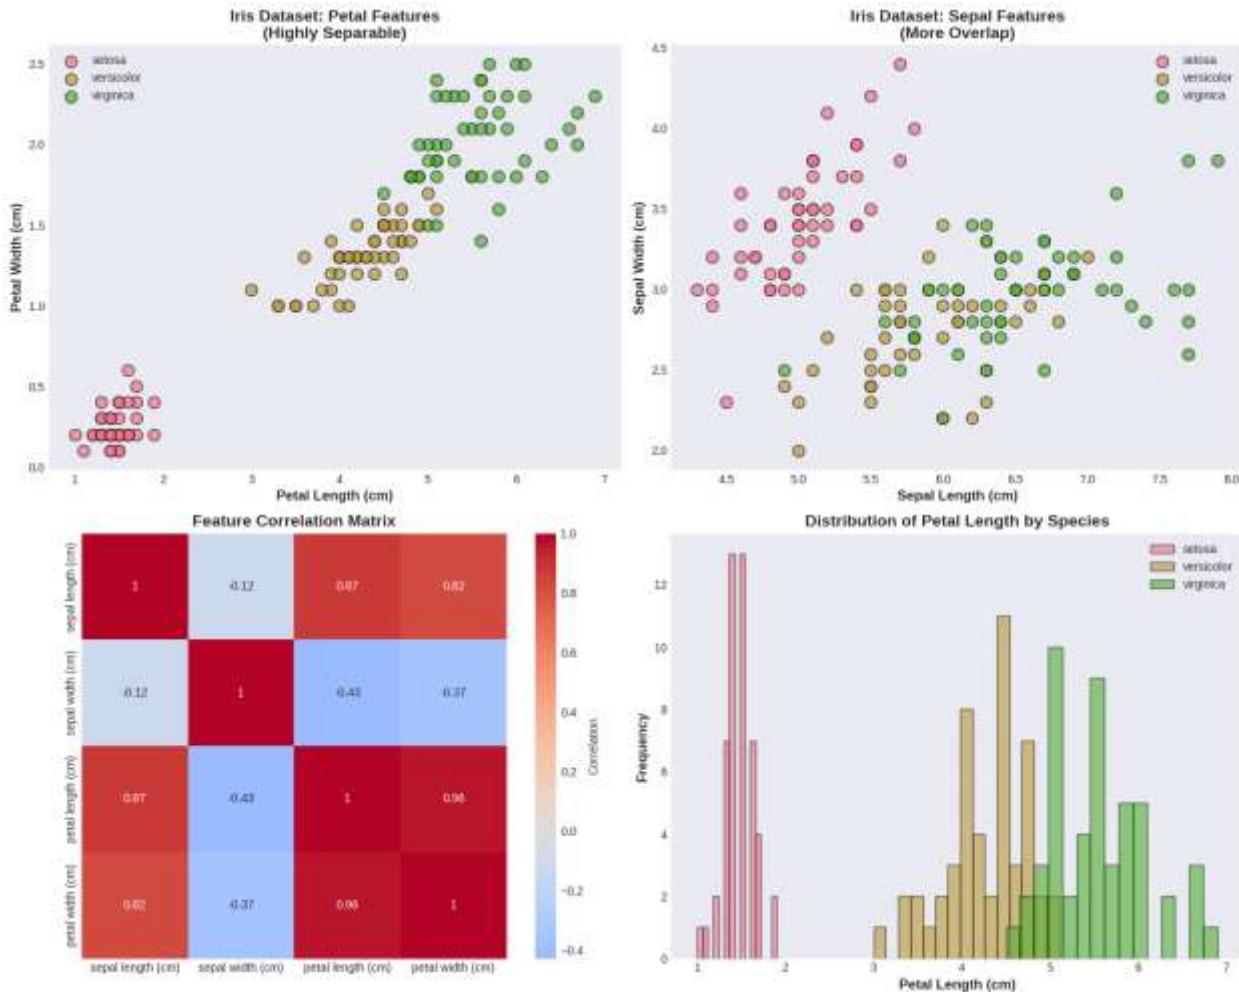


Figure 3: Exploring the Iris dataset features and distributions

As seen in Figure 3, Setosa is a clearly separated from the other two species, but Versicolor and Virginica overlap and this is clearly seen in the graph. This makes it a great test case for different kernels!

Let's understand the Kernel now:

Kernel #1: Linear Kernel

What It Does

The Linear kernel is the simplest option. It creates **straight-line decision boundaries** (or flat planes in higher dimensions).

Mathematically, it's just the dot product of two vectors:

$$K(x, y) = x^T \cdot y$$

When Should You Use It?

- The Linear kernel is your go-to choice when:

- Data is **linearly separable** (or close to it)
- Data has **many features** (high-dimensional data)
- For **fast training and prediction**
- For an **interpretable model**

What I Found

Running a Linear SVM on the Iris dataset gave me:

- **Training accuracy:** 98.10%
- **Testing accuracy:** 91.11%
- **Support vectors used:** 22 out of 105 training samples

```
=====
KERNEL 1: LINEAR KERNEL
=====
Training Linear SVM...

✓ Linear SVM trained successfully!
Training accuracy: 0.9810
Testing accuracy: 0.9111
Number of support vectors: 22
Support vectors per class: [ 2 11  9]
```

The results show that even a simple straight line does a pretty good job! Setosa was perfectly classified (100% accuracy), while Versicolor and Virginica had some confusion due to their overlap.

Kernel #2: Polynomial Kernel

What It Does

The Polynomial kernel creates **curved boundaries** using polynomial functions (Schölkopf & Smola, 2002). The formula is:

$$K(x, y) = (\gamma \cdot x^T \cdot y + r)^d$$

where d is the polynomial degree (usually 2, 3, or 4).

Think of it as fitting a curve instead of a straight line, degree 2 gives you circles and ellipses, degree 3 gives you more complex curves, and so on.

When Should You Use It?

Choose Polynomial when:

- Data has **polynomial relationships** between features

- Need **moderate non-linearity**
- For estimating the **appropriate degree** from your domain knowledge

What I Found

Using degree=3, the Polynomial SVM achieved:

- **Training accuracy:** 98.10%
- **Testing accuracy:** 91.11%
- **Support vectors used:** 24

=====

KERNEL 2: POLYNOMIAL KERNEL

=====

Training Polynomial SVM (degree=3)...

✓ Polynomial SVM trained successfully!

Training accuracy: 0.9810

Testing accuracy: 0.9111

Number of support vectors: 24

Support vectors per class: [2 12 10]

Interestingly, the accuracy matched the Linear kernel, but the decision boundary shape was noticeably different. The Polynomial kernel created smooth curves that perfectly wrapped around the Versicolor class, achieving 100% recall for it (though at the cost of some Virginica misclassifications).

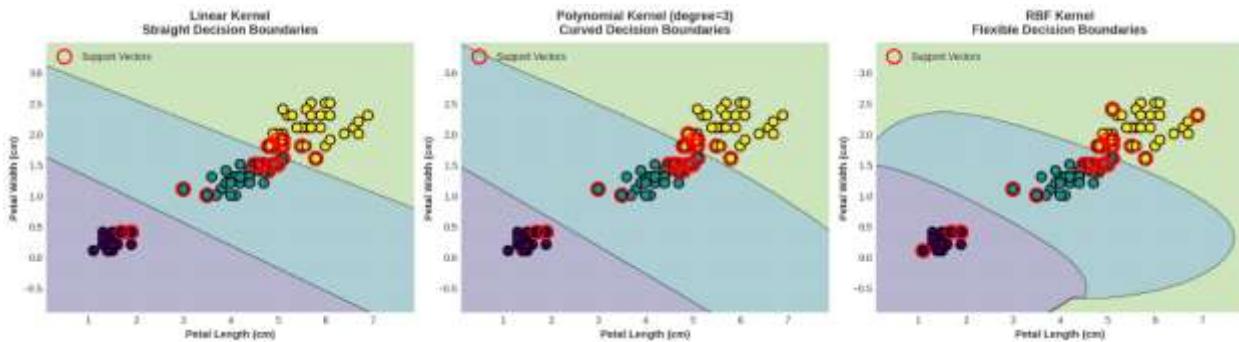


Figure 4: Comparing decision boundaries across three kernel types

Kernel #3: RBF (Radial Basis Function) Kernel

What It Does

The RBF kernel, also called the Gaussian kernel, is the **most popular and versatile** choice. Its formula is:

$$K(x, y) = \exp(-\gamma \cdot \|x - y\|^2)$$

RBF creates decision boundaries that can take almost any shape, smooth curves, tight circles around individual points, or anything in between. It essentially maps your data into infinite-dimensional space!

When Should You Use It?

RBF is your best bet when:

- **Don't know** the shape of your decision boundary
- Data has **complex, non-linear patterns**
- Want a **safe default choice** (works well in most cases)
- Willing to **tune hyperparameters** (gamma and C)

What I Found

The RBF SVM achieved:

- **Training accuracy:** 98.10%
- **Testing accuracy:** 91.11%
- **Support vectors used:** 25

```
=====
KERNEL 3: RBF (RADIAL BASIS FUNCTION) KERNEL
=====
```

```
Training RBF SVM...
```

```
✓ RBF SVM trained successfully!
Training accuracy: 0.9810
Testing accuracy: 0.9111
Number of support vectors: 25
Support vectors per class: [ 3 11 11]
```

While the accuracy matched the other kernels on this dataset, RBF's real strength shows in its flexibility. It creates smooth, organic boundaries that adapt to the local structure of the data. This makes it an actually perfect choice for complex real-world problems where we don't know the underlying pattern.

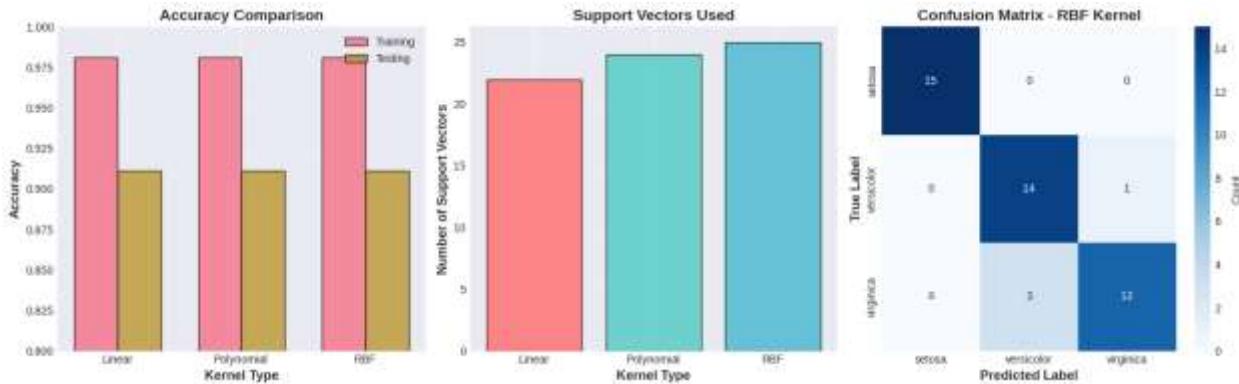


Figure 5: Comparing accuracy and support vectors used

Understanding Hyperparameters

Now that we understand what kernels are and the types of kernels. Let's move to understand the three kernels in action, and let's talk about tuning them. Two parameters control how SVMs behave: **C** and **gamma**.

The C Parameter: How Much to Care About Mistakes

C controls the trade-off between having a wide margin (simple boundary) and correctly classifying training points (complex boundary). Here are some choices for C illustrated below:

- **Small C (e.g., 0.1):** "I'm okay with some mistakes if it means a simpler, more general boundary."
- **Large C (e.g., 100):** "I want to classify every training point correctly, even if the boundary gets complex."

Think of the C parameter as how strict your SVM is about training errors. Small C is more forgiving and helps prevent overfitting.

The Gamma Parameter: How Far Should Each Point's Influence Reach?

Gamma (for RBF and Polynomial kernels) defines how much influence each training sample has. Here are some choices for choosing the gamma parameter:

- **Small gamma (e.g., 0.1):** Each point influences a **wide area**, which implies smoother, simpler boundaries
- **Large gamma (e.g., 10):** Each point only influences **nearby areas**, which implies complex, wiggly boundaries

Large gamma can lead to overfitting because the model creates tight islands around individual training points instead of learning general patterns (Schölkopf et al., 1998).

What My Experiments Showed

As C increased number of support vectors dropped from 82 to just 12, showing model was creating tighter, more complex boundaries. However, testing accuracy stayed consistent at 91.11%, suggesting that for this dataset simpler is just as good.

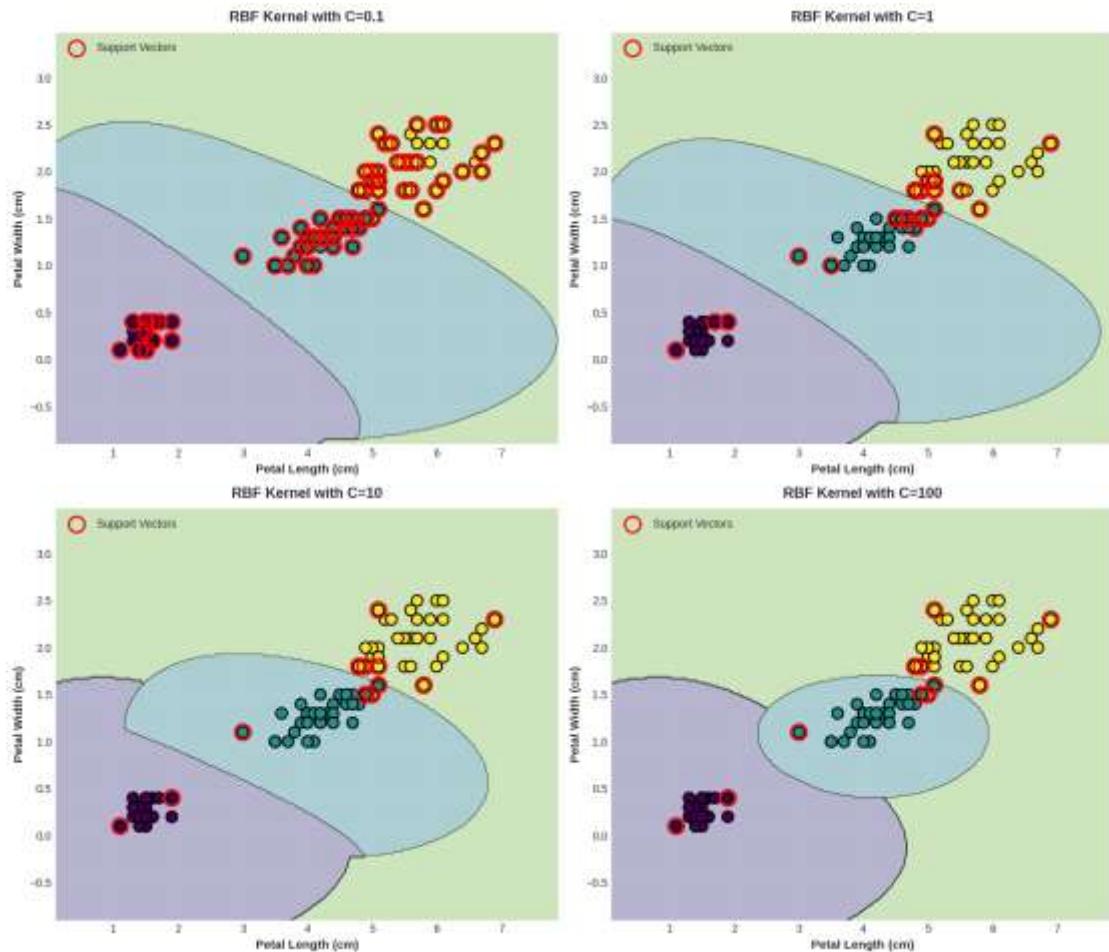


Figure 6: How C parameter changes decision boundaries

Moreover, with $\gamma=0.1$, the boundary looked almost linear (smooth and simple). With $\gamma=100$, the model created tiny islands around each training point, classic overfitting. The sweet spot was $\gamma=10$, which achieved the best testing accuracy of 95.56%.

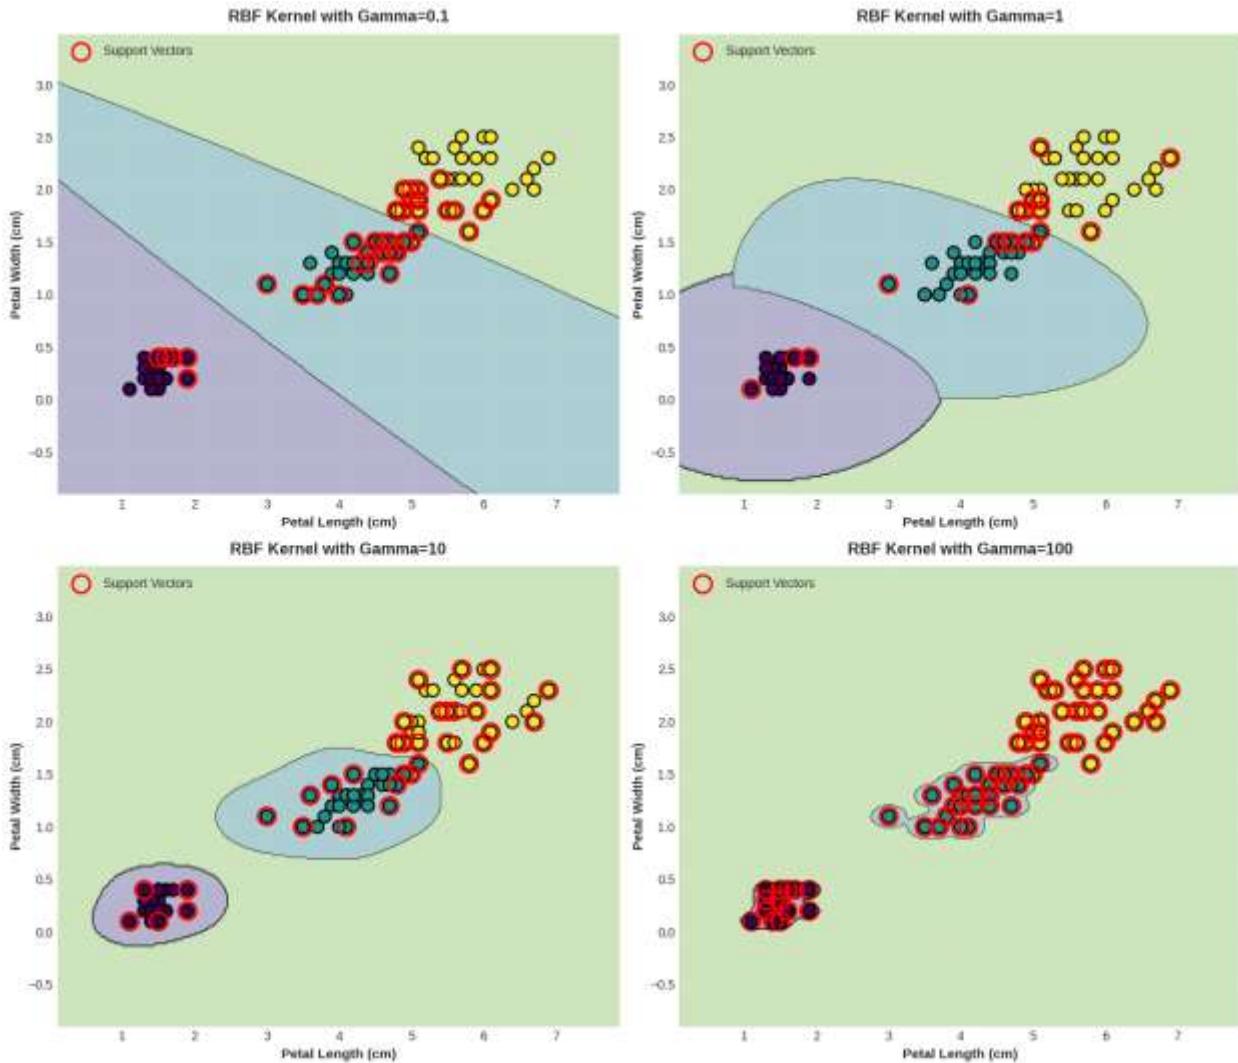


Figure 7: How gamma parameter changes boundary complexity

Finding the Best Parameters Automatically

Manually testing the parameter combinations is a difficult job, and that's where Grid Search with Cross-Validation comes in.

We tested 60 different combinations of:

- Kernels: Linear, Polynomial, RBF
- C values: 0.1, 1, 10, 100
- Gamma values: 0.001, 0.01, 0.1, 1, 10

The results showed RBF kernel with $C=0.1$ and $\gamma=0.001$ achieving 98.10% cross-validation accuracy. But here's an important insight to note that multiple combinations performed equally well. This tells us that Iris dataset is relatively easy to classify, and SVM is robust enough that several configurations work fine.

Practical Tips: What I Learned

After running all these experiments, here is advice for using SVM kernels:

1. **Start Simple, Then Add Complexity:** Always start with a Linear kernel. If it works well, great, you have a fast, interpretable model! If not, try RBF next. Only use Polynomial if you have reason to believe your data has polynomial relationships.
2. **Always Scale Your Features:** This is critical! I used StandardScaler to normalize all features to zero mean and unit variance. Without scaling, features with larger numbers will dominate, and your SVM will perform poorly. Don't skip this step!
3. **Use Cross-Validation:** A single train-test split can be misleading. Use 5-fold or 10-fold cross-validation to get a more reliable estimate of how your model will perform on new data.
4. **Watch for Overfitting:** If your training accuracy is 100% but testing accuracy is much lower, you're overfitting. Try:
 - Reducing C (more regularization):
 - Reducing gamma for RBF kernel
 - Switching to a simpler kernel
5. **When in Doubt, Use RBF:** The RBF kernel is called universal kernel for a reason; it works well in most situations. Start with default parameters ($C=1$, $\gamma='scale'$), then tune if needed.

Real-World Applications

SVMs with different kernels are used everywhere in machine learning:

- **Linear kernels:** Text classification, spam detection (high-dimensional sparse data)

- **RBF kernels:** Image recognition, bioinformatics, handwriting recognition
- **Polynomial kernels:** Computer vision tasks with known geometric relationships

The key is understanding your data and choosing accordingly.

Conclusion: Your Kernel Decision Guide

Here's a quick decision tree for choosing SVM kernels:

1. Try Linear first because it is fast, simple, and interpretable
2. If Linear doesn't work well, then try RBF because it is the most versatile.
3. If you know your data has polynomial patterns, then try Polynomial.
4. Always use Grid Search to find optimal C and gamma.
5. Never forget to scale your features.

The beauty of SVMs is that the same algorithm can handle vastly different problems just by changing the kernel. Understanding how each kernel transforms your data is the key to using them effectively.

References

- Al-Behadili, H., Grumpe, A., Dopp, C. & Wöhler, C. (2015) *Kernel trick. By transforming the original space (left) into a space of increased dimension (right) the two classes “circle” and “square” become linearly separable* [fig. 1]. Non-linear distance based large scale data classifications. ResearchGate. Available at: https://www.researchgate.net/figure/Kernel-trick-By-transforming-the-original-space-left-into-a-space-of-increased_fig1_305284381
- Cortes, C., & Vapnik, V. (1995). Support-vector networks. *Machine Learning*, 20(3), 273-297.
- Cristianini, N., & Shawe-Taylor, J. (2000). *An Introduction to Support Vector Machines and Other Kernel-based Learning Methods*. Cambridge University Press.
- curran, 2025. The Iris Dataset [online]. Available at: <https://gist.github.com/curran/a08a1080b88344b0c8a7>
- GeeksforGeeks, 2025. Support Vector Machine (SVM) Algorithm [online]. Available at: <https://www.geeksforgeeks.org/machine-learning/support-vector-machine-algorithm/>
- Pedregosa, F., Varoquaux, G., Gramfort, A., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.
- Schölkopf, B., Smola, A., & Müller, K. R. (1998). Nonlinear component analysis as a kernel eigenvalue problem. *Neural Computation*, 10(5), 1299-1319.
- Schölkopf, B., & Smola, A. J. (2002). *Learning with Kernels: Support Vector Machines, Regularization, Optimization, and Beyond*. MIT Press.
- Tech-AI-Math (2023) ‘Support Vector Machines (SVM) In Depth: Part 1: Functional Margin, Hinge Loss, Dual Problem, Lagrange Multipliers’, *Artificial Intelligence in Plain English*. Available at: <https://ai.plainenglish.io/support-vector-machines-svm-in-depth-part-1-882448c8310c>